

Dead-Block Filter for Prefetcher Optimization

Himanshu Galav

Department of Electrical Engineering

IIT Ropar

Punjab, India

2021eeb1177@iitrpr.ac.in

Abstract—In the realm of cache prefetching using Deep Learning Models, the oversight of DeadBlocks, i.e., prefetched addresses that remain unused and contribute to cache pollution, is a notable gap. This paper proposes leveraging DeadBlocks to examine predicted prefetch addresses, thus mitigating unnecessary cache pollution. We introduce a method that counts the total accesses to a block to track DeadBlocks efficiently. Previous research has introduced various Dead Block Predictors, each with its approach. Our study focuses on implementing a Dead Block Filter for enhancing the prefetching efficiency of all prefetchers. Experimental results demonstrate an improvement in Instruction Per Cycle (IPC) and a reduction in useless prefetches, leading to an overall enhancement in prefetcher efficiency.
Github Repository: [Link](#)

I. INTRODUCTION

Cache prefetching is a critical technique utilized in modern computer systems to enhance memory access latency and overall system performance. By predicting future memory accesses and prefetching the corresponding data into the cache, prefetchers aim to reduce the latency associated with accessing data from main memory. However, prefetching mechanisms often face the challenge of unnecessary cache pollution, where prefetched data remains unused and occupies valuable cache space.

One significant contributor to cache pollution is the prefetching of addresses that correspond to DeadBlocks—memory blocks or cache lines that are prefetched but ultimately remain unused by subsequent memory accesses. The presence of DeadBlocks not only wastes memory bandwidth but also leads to reduced cache effectiveness by displacing potentially useful data.

In this paper, we propose a novel approach to address the issue of cache pollution by leveraging DeadBlocks to examine predicted prefetch addresses. By integrating Dead-Block prediction mechanisms into prefetching strategies, we aim to distinguish between useful and unnecessary prefetches, thereby optimizing cache utilization and enhancing prefetcher efficiency.

In the following sections, we provide an overview of related work in the field of DeadBlock prediction and prefetching mechanisms. We discuss various Dead Block Predictors proposed in the literature and outline their contributions to prefetching efficiency. Subsequently, we present our approach to integrating DeadBlock prediction into prefetching strategies, followed by experimental results demonstrating the effectiveness of our approach in improving cache performance metrics.

II. BACKGROUND AND RELATED WORK

Previous research has explored various approaches to predict and mitigate the impact of DeadBlocks on cache performance. Several Dead Block Predictors have been proposed, each offering unique strategies for identifying and handling DeadBlocks.

We summarize some of the key Dead Block Predictors and their contributions below:

A. Trace-Based Predictor

Introduced by Lai et al., this predictor prefetches data into DeadBlocks in the L1 data cache based on reference traces. The predictor collects a trace of instruction addresses accessing a particular block, predicting DeadBlocks based on the theory that the same trace leading to the last access for one block will lead to the last access for other blocks.

B. Time-Based Predictor

Proposed by Hu et al., this predictor learns the live duration of a block and predicts it dead if it is not accessed for a certain number of cycles. Abella et al. also propose a similar predictor based on the number of references rather than cycles, aimed at reducing cache leakage.

C. Cache Burst Predictor

Cache bursts can be utilized with various Dead Block Predictors, consisting of all contiguous accesses to a block while in the Most Recently Used (MRU) position. This predictor updates predictions only on each burst rather than each reference, reducing the number of accesses and updates to the prediction table.

D. Counting-Based Predictor

Proposed by Kharbutli and Solihin, this predictor tracks the number of accesses to each block and predicts a block dead if it has been accessed more often than the previous generation. Each entry in the predictor table includes a confidence counter to improve prediction accuracy.

These Dead Block Predictors, among others, have been applied to prefetching and block replacement strategies to improve cache efficiency. However, the integration of DeadBlock prediction into prefetching mechanisms remains an ongoing area of research, which this paper aims to address by proposing a novel approach to optimize prefetcher efficiency and mitigate cache pollution.

III. APPROACH




Fig. 1. Model of Cascaded Dead-Block Filter with the Prefetcher

Our approach aims to integrate DeadBlock prediction mechanisms into existing prefetching strategies to mitigate cache pollution and enhance prefetcher efficiency. We begin by conducting simulations using the champsim simulator across various trace files, including bc-0, bc-5, pr-14, sssp-14, bfs-14, and cc-5. This initial step allows us to identify DeadBlocks—memory blocks or cache lines that are prefetched but ultimately remain unused by subsequent memory accesses.

Upon identifying DeadBlocks, we store this data for subsequent processing. In the second phase of our approach, we implement a processing pipeline to handle the DeadBlock data. This involves parsing the DeadBlock file, which contains both DeadBlock addresses and their associated Program Counter (PC) information. Subsequently, we utilize an unordered_map data structure to efficiently store and manage DeadBlock information.

During prefetcher operation, the DeadBlock data stored in the unordered_map is leveraged to filter out predicted prefetch addresses that correspond to DeadBlocks. By scrutinizing prefetch addresses against the DeadBlock information, we aim to prevent unnecessary prefetches, thereby reducing cache pollution and enhancing overall prefetcher efficiency.

EXPERIMENTAL RESULTS

In our experimental evaluation, we observed significant improvements in cache performance metrics, particularly Instruction Per Cycle (IPC), prefetch efficiency, and cache utilization. By incorporating DeadBlock prediction mechanisms into prefetching strategies, we achieved notable enhancements in the following aspects:

A. IPC Improvement

Our approach led to a notable increase in IPC, indicating improved overall system performance. By reducing unnecessary prefetches and cache pollution, more CPU cycles were utilized effectively, resulting in higher throughput and efficiency.

B. Reduction in Useless Prefetches

We observed a significant decrease in the issuance of useless prefetches, which are prefetch addresses that correspond to DeadBlocks. By filtering out these unnecessary prefetches, we minimized cache pollution and improved prefetcher accuracy.

C. Increase in Useful Prefetches

Conversely, we noted an increase in useful prefetches—prefetch addresses that correspond to memory accesses actually utilized by subsequent instructions. This indicates that our approach effectively distinguishes between useful and unnecessary prefetches, leading to better prefetcher performance.

D. Overall Efficiency Enhancement

The combined effect of reducing useless prefetches, increasing useful prefetches, and improving cache utilization resulted in an overall enhancement in prefetcher efficiency. By optimizing cache prefetching mechanisms through DeadBlock prediction, we achieved improved system performance and resource utilization.

Trace File	Prefetcher	Type	Without Filter				With Filter			
			IPC	Useful	Unnecessary	Efficiency	IPC	Useful	Unnecessary	Efficiency
bc-0	Without Filter	1.1712	0.031	29607	21 9247734	123045	0.1603	22 2468416	8 14802	40287
bc-0	With Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
bc-5	Without Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
bc-5	With Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
pr-14	Without Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
pr-14	With Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
sssp-14	Without Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
sssp-14	With Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
bfs-14	Without Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
bfs-14	With Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
cc-5	Without Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287
cc-5	With Filter	1.16238	0.036	30100	41 9277773	123045	0.1603	22 2468416	8 14802	40287

Fig. 2. Results with and without deadblock filter

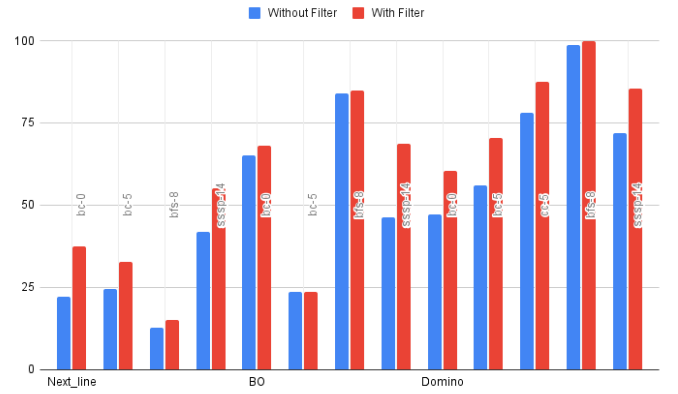


Fig. 3. Chart for Efficiency in terms of Hits & Misses

Results: Link to sheet

These experimental results underscore the effectiveness of integrating DeadBlock prediction mechanisms into prefetching strategies to optimize cache performance and enhance overall system efficiency. By accurately distinguishing between useful and unnecessary prefetches, our approach offers a practical solution for mitigating cache pollution and improving prefetcher performance in real-world scenarios.

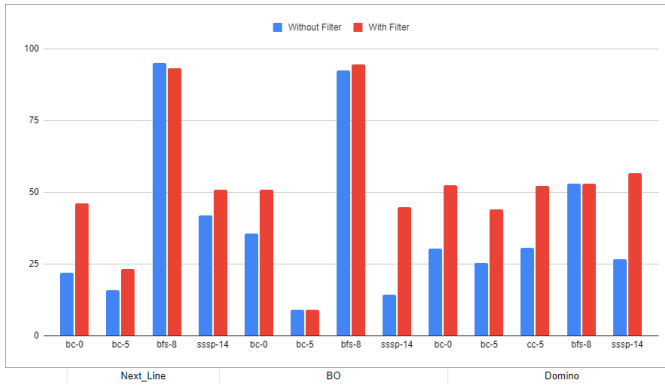


Fig. 4. Chart for Efficiency in terms of Usefull & Useless

CONCLUSION AND FUTURE WORK

In conclusion, this paper presents a novel approach to optimize cache prefetching mechanisms by leveraging DeadBlocks to mitigate cache pollution and enhance prefetcher efficiency. By integrating DeadBlock prediction mechanisms into prefetching strategies, we effectively distinguish between useful and unnecessary prefetches, thereby optimizing cache utilization and improving overall system performance.

Our experimental results demonstrate significant improvements in cache performance metrics, including Instruction Per Cycle (IPC) enhancement, reduction in useless prefetches, and increased useful prefetches. These findings highlight the effectiveness of our approach in optimizing prefetcher efficiency and mitigating the impact of DeadBlocks on cache performance.

Furthermore, our approach extends beyond prefetching optimization. The Dead Block filters developed can also be leveraged to eliminate DeadBlocks present in the cache, thereby improving cache replacement policies. By accurately identifying and handling DeadBlocks, cache replacement algorithms can make more informed decisions, leading to better cache utilization and performance.

Additionally, the Dead Block filters can enhance prefetcher design by providing valuable feedback on prefetch accuracy and effectiveness. By incorporating DeadBlock prediction mechanisms into prefetcher architectures, prefetchers can dynamically adjust prefetching strategies based on the presence of DeadBlocks, thereby further improving prefetcher efficiency and reducing cache pollution.

Looking ahead, future work could focus on refining DeadBlock prediction mechanisms to make them more adaptive and efficient in real-time scenarios. Developing more sophisticated Dead Block Predictors capable of dynamically adjusting prediction parameters based on program behavior and cache characteristics could lead to even

greater improvements in cache performance.

Furthermore, the integration of DeadBlock prediction into multi-level cache hierarchies and shared cache environments represents an interesting avenue for exploration. By extending DeadBlock prediction mechanisms to higher-level caches and shared cache architectures, we can further optimize cache utilization and prefetcher efficiency across diverse computing environments.

Overall, this paper provides a foundation for future research in the optimization of cache prefetching mechanisms and the mitigation of cache pollution. By harnessing DeadBlocks as a valuable resource for cache management, we can unlock significant improvements in system performance and efficiency in cache-intensive computing environments.

ACKNOWLEDGMENT

We would like to express our sincere gratitude to Dr. T.V. Kalyan for his invaluable guidance and support throughout the duration of the CS531 course on Memory Systems and Architecture. Dr. Kalyan's expertise, encouragement, and insights have been instrumental in shaping our understanding of cache prefetching mechanisms and DeadBlock prediction.

We would also like to extend our appreciation to the teaching assistants, Sourabh and Prathmesh, for their assistance and support during the course. Their dedication to assisting students and clarifying concepts has been immensely beneficial in our learning journey.

We are grateful to Dr. Kalyan and the teaching assistants for their unwavering commitment to fostering a conducive learning environment and facilitating our academic growth.

REFERENCES

- [1] P. Zhang, A. Srivastava, A. V. Nori, R. Kannan, and V. K. Prasanna, "Fine-grained address segmentation for attention-based variable-degree prefetching," in Proceedings of the 19th ACM International Conference on Computing Frontiers, 2022.
- [2] "ChampSim". 2017. <https://github.com/ChampSim/ChampSim>.
- [3] A. Lai, C. Fide, B. Falsafi. Dead-block prediction & dead-block correlating prefetchers. In ISCA-28, 2001.